



OPTIMIZING SQL QUERIES

TABLE OF CONTENTS

MISSIONS TO COMPLETE.....	4
EXERCISE TRAFFIC ANALOGY.....	4
MISSION TOOL.....	5
EXECUTION PLAN TOOL.....	5
CLASS EXERCISE GENERATING AND INTERPRETING.....	6
INDEPENDENT EXERCISE.....	9
MISSION 1 REVISIT.....	10
MISSION MAP.....	11
INDEXES.....	11
LOCATING INDEXES.....	12
CLUSTERED VS. NON-CLUSTERED.....	13
INDEXED COLUMNS.....	13
CLASS EXERCISE UTILIZING INDEXES.....	14
INDEPENDENT EXERCISE.....	15
MISSION 2 REVISIT.....	15
MISSION BRIDGE.....	17
JOINS.....	17
JOINING AT LEAST FIVE TABLES	18
TABLE JOIN RELATIONSHIP	19
INDEPENDENT EXERCISE.....	20
MISSION 3 REVISIT.....	20
CLOSING EXERCISE.....	21
COMPLETED MISSIONS.....	24
ADDITIONAL RESOURCES.....	25
SQL Sentry Plan Explorer Tool	25

INTERNAL USE ONLY

OPTIMIZING SQL QUERIES

OVERVIEW

As a Developer who often writes SQL code, continuously learning to optimize your SQL queries is crucial when working in the Shore 1 Server. This hands-on training will focus on the use of a query optimization tool within SQL Server Management Studio (SSMS), the use of indexes, and the use of JOINs with 5 or more tables. Learning to optimize your SQL queries results in faster running queries, which is a quicker return on data. Additionally, it will reduce server slowness and potential risk of the server crashing.



INTERNAL USE ONLY

MISSIONS TO COMPLETE


Accomplish these three missions to save the city (Shore 1 Server) from ticking time bombs (inefficient queries)!

- Practice generating and interpreting execution plans for several SQL queries.
- Practice optimizing several SQL queries using indexes.
- Practice optimizing several SQL queries through proficient use of JOINs with 5 or more tables at a time.



EXERCISE | TRAFFIC ANALOGY

After watching the second video shown by your trainer, use the space provided to note anything you see that is causing the slowness and traffic jams, as well as ways that it can be avoided.



MISSION TOOL

The SSMS has a tool that is crucial to understanding and optimizing SQL queries – Execution Plan tool. When enabled, this tool allows you to generate an execution plan for your query, and the plan shows the overall performance of your query. It also breaks your query down to point out areas where your query could be more efficient, so that you can modify your query based on the findings. This tool is required to be used with all your written queries.

EXECUTION PLAN TOOL


The Execution Plan tool is designed to help you understand the steps taken by SQL Server to execute a query. It provides a visual representation of the query's execution plan, showing the flow of data and the operations performed. This tool is essential for optimizing queries and improving their performance.

- “GPS Navigation System” for query.
- Generates an execution plan.
- Outlines key information about the query.
 - Cost of operations.
 - Amount of data processed.



INTERNAL USE ONLY


NOTES:



CLASS EXERCISE | GENERATING AND INTERPRETING

Once the SQL query is written, you can use the Execution Plan tool to generate an execution plan for that query by turning the tool on and then executing your query. The guidance below shows you how to enable this tool and how to read the execution plan.

1. Open the SSMS and connect to the following server: **S01VDSHR01AGL**.
2. Click **New Query**.



3. Copy and paste the following SQL query into the new query window.

```
SET STATISTICS IO ON
SET STATISTICS TIME ON


SELECT TOP 1000 * from [dbo].[loan_conditions]
WHERE loanrecordid = 21534067
ORDER BY categoryid

SET STATISTICS IO OFF
SET STATISTICS TIME OFF
```


NOTE: Including **SET STATISTICS IO/TIME** in your query makes it so that you can monitor the CPU time and number of Logical Reads when you're query runs. These two metrics can be found within the "Messages" tab.

4. Enable the tool in the SSMS.

- To enable the Execution Plan tool, right-click anywhere in the query window and select **Include Actual Execution Plan**.



NOTE: Once you have enabled the Execution Plan tool, execute your written query. A new tab for **Execution plan** will display in the bottom window.




The screenshot shows the SSMS interface with the 'Execution plan' tab selected in the bottom navigation bar. The main pane displays the execution plan for a query. The plan shows a 'Clustered Index Seek' operation on the [loan_conditions].[PK_loan_...] index, costing 100% of the query. The 'Results' tab shows the query results below the plan.

```

  Results Messages Execution plan
Query 1: Query cost (relative to the batch): 100%
select top 1000 * from [dbo].[loan_conditions] where lenderdatabaseid = 1 and loanrecordid = 21534067 order by categoryid
  
```

SELECT	Top	Clustered Index Seek (Cluste...
Cost: 0 %	Cost: 0 %	[loan_conditions].[PK_loan_...]
0.000s	0.000s	
18 of	18 of	
90 (20%)	90 (20%)	

EXAMPLE: When reviewing an execution plan, focus on these key areas, including number of rows being processed.



INTERNAL USE ONLY

Key Lookup (Clustered)

Uses a supplied clustering key to lookup on a table that has a clustered index.

Physical Operation	Key Lookup
Logical Operation	Key Lookup
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows for All Executions	18
Actual Number of Rows Read	18
Actual Number of Batches	0
Estimated Operator Cost	3.77351 (99%)

This number should make sense with what the query is asking for.

INDEPENDENT EXERCISE

Copy and paste the following SQL query into the SSMS. Generate and interpret the execution plan, focusing on the operator types (i.e., Index Seek, Key Lookup), cost, CPU time, and logical reads. Use the space provided to summarize the differences between this query and the previous query.

```
SET STATISTICS IO ON  
SET STATISTICS TIME ON  
  
SELECT TOP 1000 * FROM [dbo].[loan_conditions]  
WHERE lenderdatabaseid = 1 AND loanrecordid = 21534067  
ORDER BY categoryid  
  
SET STATISTICS IO OFF  
SET STATISTICS TIME OFF
```



DIFFERENCES:



MISSION 1 | REVISIT

Complete the following questions related to execution plans to allow yourself to critically think about what we just learned.

How can you tell if a query is using an index when interpreting an execution plan?

If you see a Key Lookup or Table Scan in your execution plan, what does that suggest about your query?

How does an execution plan help to identify bottlenecks in real-time?

What are 2-3 metrics you should focus on when reviewing execution plans and the Messages tab?

What does an operation with a high cost suggest about your query?

INTERNAL USE ONLY


MISSION MAP

Similarly to a detailed map of the "Shore 1 Server" city, an index is a way to quickly find and return data from a table without the server having to sort through every row in the table every time a database table is accessed. Indexes should always be reviewed prior to optimizing your SQL query to identify what indexes are available and what columns of the table they cover. It is not always possible to use an index, but you must always try.

INDEXES

Indexes speed up the return on data and it's a cheaper operator to use compared to Table Scans and Key Lookups.

- Always utilize indexes, when possible.
- Speeds up data retrieval, quicker return on data.
- Reduce server load, less data to examine.
- SQL server selects index, based on written query.




NOTES:



LOCATING INDEXES

To be able to optimize your queries so the SQL server likely will utilize an index, we must know where to find the index for a table and what columns of the table that the index covers. Follow the guidance below to do this.

1. Click on the necessary database (**i.e., Integra_prod (Synchronized)**).
2. Click on the necessary table you're using (**i.e., BRRWER.DriversLicense**).
3. Click on the **Indexes folder** to expand the available index list.



NOTE: Not all tables have indexes created. If one doesn't exist and you think the floor could benefit from having one, reach out to the DBA and DEA teams to create one, but always include your Team Lead in the communication.

CLUSTERED VS. NON-CLUSTERED


There are clustered and non-clustered indexes. Clustered indexes are faster than non-clustered indexes. See the table below for the differences.

CLUSTERED VS NON-CLUSTERED INDEXES		
FEATURE	CLUSTERED	NON-CLUSTERED
Physical Order	Affects physical order of data in table.	Physical data unchanged in table.
Number per Table	One	Multiple
Data Retrieval	Faster	Slower
Best For	Range queries	Specific column queries

INDEXED COLUMNS

Knowing which **key columns**, or “keys”, are covered by an index improves your query planning. You cannot modify your query in a way that the SQL server is likely to use that index unless you know what columns are indexed. Double-click on a specific index to see what columns are covered.

NOTE: For the SQL Server to utilize an index, include all indexed columns or minimally, the top column(s) in your query criteria. For instance, if using the below index, incorporate all four columns or at least the BorrowerID. The BorrowerID can't be skipped as it's the top listed column.



CLASS EXERCISE | UTILIZING INDEXES

We talked about the importance of modifying queries in a way that they use available indexes to increase query optimization. Let's look at the difference in efficiency between using indexes and using Table Scans and Key Lookups. Follow along with the trainer as you generate an execution plan on two SQL queries shown below. As a class, make note of the differences when using an index vs. not using an index.

--De-optimized:

```
SET STATISTICS IO ON
SET STATISTICS TIME ON

SELECT *
FROM dbo.USFS_OrgChart uoc WITH(NOLOCK)
WHERE uoc.userid = 22316

SET STATISTICS IO OFF
SET STATISTICS TIME OFF
```

--Optimized:

```
SET STATISTICS IO ON
SET STATISTICS TIME ON

SELECT o.FullName
    ,o.JobTitle
    ,o.manager_FullName
FROM dbo.USFS_OrgChart o WITH(NOLOCK)
WHERE o.domainname = 'jhalim'
    AND o.userid = 22316

SET STATISTICS IO OFF
SET STATISTICS TIME OFF
```

DIFFERENCES:

INDEPENDENT EXERCISE


Generate and interpret an execution plan for the following query, then optimize the query. Remember to review the available indexes prior to optimizing the query.

--de-optimized:

```
SET STATISTICS IO ON  
SET STATISTICS TIME ON  
  
SELECT re.[name]  
      ,re.alias2  
      ,re.entityid  
FROM dbo.rolodex_entity re WITH(NOLOCK)  
WHERE re.entityid = 89;  
  
SET STATISTICS IO OFF  
SET STATISTICS TIME OFF
```




OPTIMIZED QUERY:



MISSION 2 | REVISIT

Think about what you've learned and your biggest challenge(s) when using indexes. Using the space provided, write down what your biggest challenge will be with using indexes when you go back to your desk and why.



** INTERNAL USE ONLY **

MISSION BRIDGE

JOINS act as the bridges in Shore 1 Server city. Bridges in the city connect us to different areas, allowing us to move from one side of the city to the other more quickly to diffuse the ticking time bombs! JOINS serve a similar purpose in SQL by allowing us to combine data from separate tables to be accessed together, making data retrieval faster and more efficient on the server.

JOINS

Below are key points to keep in mind when JOINing five or more tables.

- Order by smallest, most filtered result sets first.
 - This is not always most optimal.
- Concept of JOIN type remains the same.
- Use INNER JOIN when possible.
- JOIN on indexed columns. (Order matters).



NOTE: JOINS used in conjunction with a Function must be reviewed by sr. team member or team lead.

EXAMPLE: Order JOINs by smallest, most filtered result sets first. The below example shows this by starting with selecting all the borrowers with a credit score over 700 first. Then, it JOINs the smaller result set of borrowers to the Loans table (since it has the least number of rows), and then the Payments table to the smaller set of loans. This way, the server is filtering first and then JOINing smaller sets of data.

Tables Find Credit Scores Above 700	
Loans	1,000 Rows
Borrowers	10,000 Rows
Payments	100,000 Rows

```

SELECT *
FROM Borrowers
JOIN Loans ON Borrowers.borrower_id = Loans.borrower_id
JOIN Payments ON Loans.loan_id = Payments.loan_id
WHERE Borrowers.credit_score > 700;
  
```

JOINING AT LEAST FIVE TABLES

We talked about the importance of JOINing on indexed key columns, ensuring the key columns are listed in the order in which they appear in the indexes, and ensuring not to use any tables that are not necessary. Let's get practice as a class with identifying these things in a query and modifying the query to be more optimized.

--De-Optimized:

--Loans Currently in Submission status. We want to grab the Loan number, Borrowers name and the Senior UW on the Loan.

```

SELECT lm.loanid
      ,CONCAT(cm.firstname, ' ', cm.lastname) BorrowerName
      ,CONCAT(rc.firstname, ' ', rc.lastname) srUWName
FROM dbo.customer_main cm WITH (NOLOCK)
  INNER JOIN dbo.customer_group cg WITH (NOLOCK)
    ON cg.customerrecordid = cm.customerrecordid
    AND cg.lenderdatabaseid = cm.lenderdatabaseid
  INNER JOIN dbo.loan_main lm WITH (NOLOCK)
    ON lm.customergroupid = cg.customergroupid
    AND lm.lenderdatabaseid = cg.lenderdatabaseid
  INNER JOIN dbo.loan_status ls WITH (NOLOCK)
    ON ls.lenderdatabaseid = lm.lenderdatabaseid
    AND ls.statusid = lm.statusid
    AND ls.loanrecordid = lm.loanrecordid
  INNER JOIN dbo.loan_channelcontacts lc WITH (NOLOCK)
    ON lc.loanrecordid = ls.loanrecordid
    AND lc.lenderdatabaseid = LS.LENDERDATABASEID
    AND lc.contactcategoryid = 8 --senior underwriter
  INNER JOIN dbo.rolodex_contacts rc WITH (NOLOCK)
    ON rc.contactid = lc.contactid
    AND rc.entityid = lc.entityid
  INNER JOIN dbo.usfs_testloanindicator T WITH (NOLOCK)
    ON t.lenderdatabaseid = lm.lenderdatabaseid
    AND t.loanrecordid = lm.loanrecordid
    AND t.testloanind = 0 --removing test loans
WHERE ls.statusid IN ( 2, 49 ) --UW status/submission
  
```

INTERNAL USE ONLY

--Optimized:

```

SELECT lm.loanid
      ,CONCAT(cm.firstname, ' ', cm.lastname) BorrowerName
      ,CONCAT(rc.firstname, ' ', rc.lastname) srUWName
FROM dbo.loan_main lm WITH (NOLOCK)
  
```


```

INNER JOIN dbo.customer_group cg WITH (NOLOCK)
    ON cg.lenderdatabaseid = lm.lenderdatabaseid
        AND cg.customergroupid = lm.customergroupid
INNER JOIN dbo.customer_main cm WITH (NOLOCK)
    ON lm.lenderdatabaseid = lm.lenderdatabaseid
        AND cg.customerrecordid = cm.customerrecordid
INNER JOIN dbo.loan_channelcontacts lc WITH (NOLOCK)
    ON lc.loanrecordid = lm.loanrecordid
        AND lc.lenderdatabaseid = lm.lenderdatabaseid
        AND lc.contactcategoryid = 8 --srUW
INNER JOIN dbo.rolodex_contacts rc WITH (NOLOCK)
    ON rc.lenderdatabaseid = lc.lenderdatabaseid
        AND rc.entityid = lc.entityid
        AND rc.contactlenderdatabaseid = lc.contactlenderdatabaseid
        AND rc.contactid = lc.contactid
WHERE lm.statusid IN ( 2, 49 );

```

TABLE JOIN RELATIONSHIP

The below diagram represents all the tables and what each table is JOINed to from the above query example. The bolded names are the keys that each table have in common.



INDEPENDENT EXERCISE


It's your turn to practice optimizing a query with multiple JOINs. Review the query below, then optimize the query.

--De-optimized:

```

SELECT lm.loanid
      ,CONCAT(cm.firstname, ' ', cm.lastname)
    BorrowerName
      ,CONCAT(rc.firstname, ' ', rc.lastname) srUWName
FROM dbo.loan_main lm WITH (NOLOCK)
  JOIN dbo.customer_group cg WITH (NOLOCK)
    ON cg.customergroupid = lm.customergroupid
  JOIN dbo.customer_main cm WITH (NOLOCK)
    ON cg.customerrecordid = cm.customerrecordid
  JOIN dbo.loan_channelcontacts lcc WITH (NOLOCK)
    ON lcc.loanrecordid = lm.loanrecordid
      AND lcc.contactcategoryid = 8 --loan officer
  JOIN dbo.rolodex_contacts rc WITH (NOLOCK)
    ON rc.contactid = lcc.contactid
  JOIN dbo.USFS_TestLoanIndicator t WITH (NOLOCK)
    ON t.LoanRecordID = lm.loanrecordid
      AND t.TestLoanIND = 0 --removing test loans
WHERE lm.statusid IN ( 2, 49 );

```



MISSION 3 | REVISIT

Take a moment to critically think about the effects of properly JOINing tables, and answer the following question using the space provided: *How will the proper use of JOINs in your queries make your queries run faster and what impact does that have on the Shore 1 server?*



CLOSING EXERCISE

Review the following queries and then optimize them. Remember to run them through the Execution Plan tool to review things like cost, operator type, CPU time, and number of logical reads. Remember to review available indexes and JOIN criteria. Copy and paste your optimized query in the spaces provided for future reference.

1. Optimize the following query:

--De-optimized:


```
SET STATISTICS IO ON
SET STATISTICS TIME ON

SELECT lcd.debtid, cd.name, cd.numberofpayments, lcd.exclusionreason, cd.debtttype,
lcd.includetpayment, lcd.includetbalance, cd.balance, cd.payment, cd.lienposition,
lcd.payoffrequired
    FROM    loan_customerdebt lcd
    JOIN   loan_main lm  ON  lm.loanrecordid = lcd.loanrecordid
    JOIN   customer_debt cd  ON cd.debtid = lcd.debtid
        AND cd.debtldid = lcd.debtldid
    WHERE  lcd.loanrecordid = 21493577

SET STATISTICS IO OFF
SET STATISTICS TIME OFF
```

INTERNAL USE ONLY

PASTE OPTIMIZED QUERY



2. Optimize the following query:

--De-optimized:

```
SET STATISTICS IO ON
SET STATISTICS TIME ON

SELECT lm.lenderdatabaseid AS LenderDatabaseID, lm.loanrecordid AS LoanRecordID,
lm.loanid AS LoanNumber, clm.ContactEmailAddress
    FROM dbo.loan_main lm
    JOIN dbo.custom_loanmain clm ON clm.loanrecordid = lm.loanrecordid AND
clm.lenderdatabaseid = lm.lenderdatabaseid

    WHERE lm.loanrecordid = 21475838

SET STATISTICS IO OFF
SET STATISTICS TIME OFF
```

PASTE OPTIMIZED QUERY

3. Optimize the following query:

```
--de-optimized

SET STATISTICS IO ON;
SET STATISTICS TIME ON;

SELECT lm.lenderdatabaseid
    ,lm.loanrecordid
    ,lm.loanid
    ,lpc.dispositiondate
    ,sls.statusdescription
    ,p.productdescription
FROM dbo.loan_main AS lm WITH (NOLOCK)
JOIN dbo.loan_postclosing AS lpc WITH (NOLOCK)
    ON lpc.loanrecordid = lm.loanrecordid
        AND lpc.disposition = 1
JOIN dbo.property_main AS pm WITH (NOLOCK)
    ON pm.propertyrecordid = lm.propertyrecordid
        AND pm.state = 'ND'
JOIN dbo.loan_huditem AS lhi WITH (NOLOCK)
    ON lhi.loanrecordid = lm.loanrecordid
        AND lhi.lenderdatabaseid = lm.lenderdatabaseid
        AND lhi.currenthudnumber IN ( 804, 805 )
JOIN dbo.setups_loanstatus sls WITH (NOLOCK)
    ON sls.statusid = lm.statusid
JOIN dbo.product_main p WITH (NOLOCK)
    ON p.productid = lm.productid;

SET STATISTICS IO ON;
SET STATISTICS TIME ON;
```

INTERNAL USE ONLY

PASTE OPTIMIZED QUERY

COMPLETED MISSIONS

Congratulations! You have accomplished these three missions and saved the city from all ticking time bombs!

- Practice generating and interpreting execution plans for several SQL queries.
- Practice optimizing several SQL queries using indexes.
- Practice optimizing several SQL queries through proficient use of JOINs with 5 or more tables at a time.




ADDITIONAL RESOURCES

See the below additional videos and readings that cover what was included in today's SQL training.


- <https://www.youtube.com/watch?v=SMw2knRullE>
- <https://www.youtube.com/watch?v=uqkf1VVZulg&t=474s>
- <https://learning.oreilly.com/library/view/sql-tuning/0596005733/>

SQL SENTRY PLAN EXPLORER TOOL

SQL Sentry Plan Explorer tool is another version of an execution plan that can be installed through a third-party website based upon personal preference. Some team members find it slightly easier to compare execution plans between multiple queries; however, both execution plan tools run the same way and provide the same information.



Click [here](#) to install. Once installed, turn on actual execution plan and run query in SSMS. Once query runs, right click anywhere in the execution plan and select **View with SQL Sentry Plan Explorer**.



If you have any questions or concerns,
please reach out to ittraining@uwm.com.